

# SSD Training Tutorial

J. Cory – Embedded Vision Specialist  
1.11.19

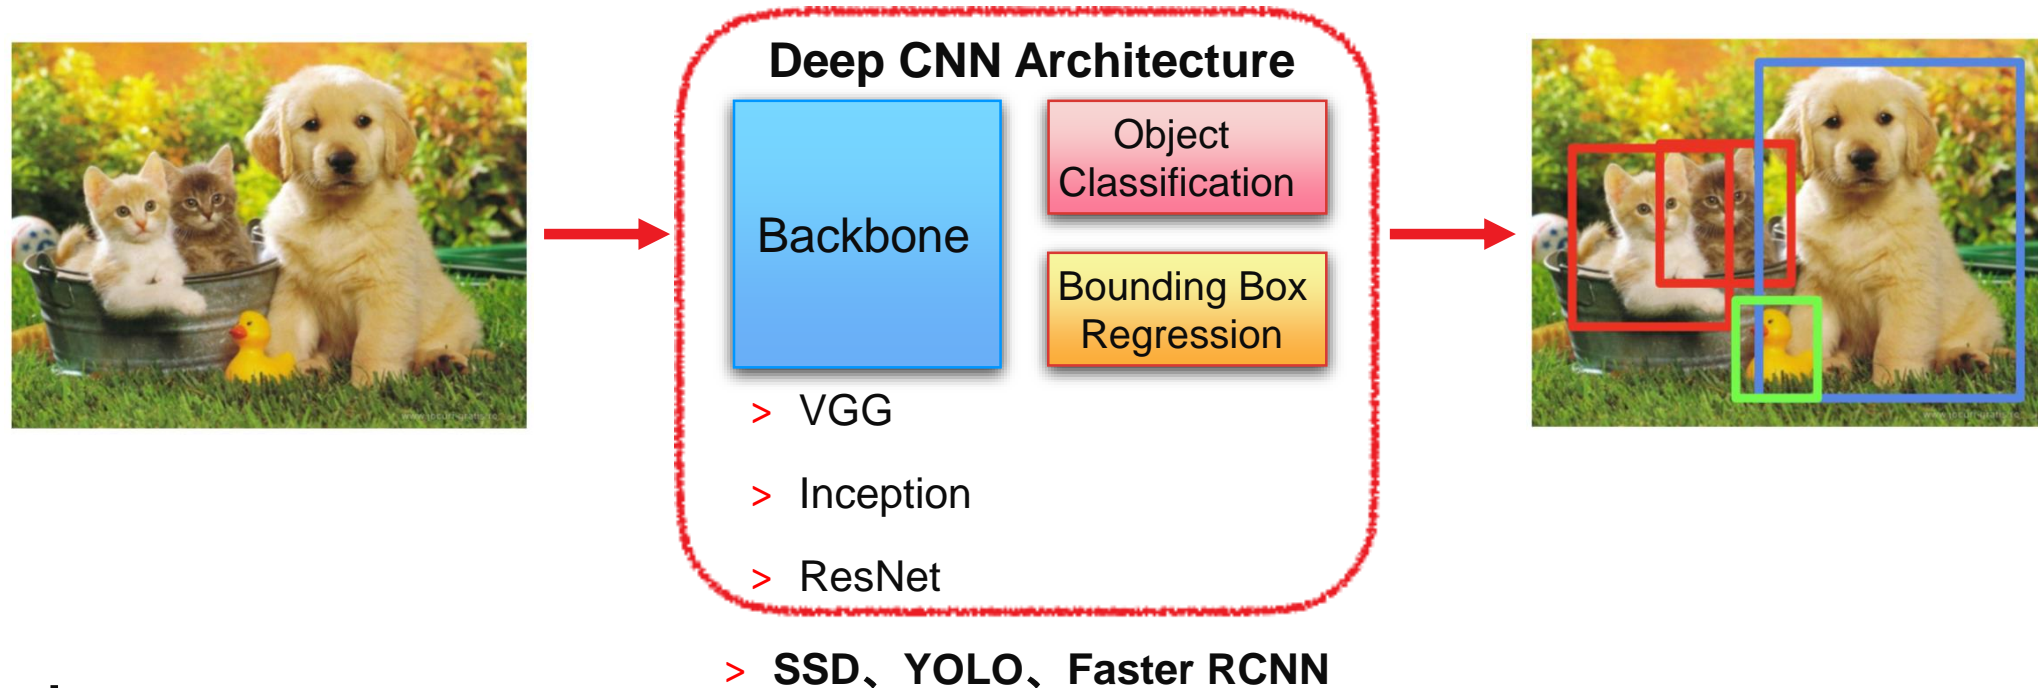
*Credit to: Louie Valena for detection metrics slides*



# Introduction



# ML Task: Detection



## > Task:

- >> Input: Image
- >> Output: Class label and location of objects
- >> Evaluation metric:
  - >> mAP, mean average precision
  - >> Precision at fixed recall
  - >> Recall at fixed precision

## > Applications:


- >> Face detection
- >> License plate detection
- >> Pedestrian, cyclist and car detection in surveillance/automotive.

# Detection Metrics




# Metrics: Precision and Recall

- > Precision: how many selected items are relevant?
- > Recall: how many relevant items are selected?
- > Ex.: 12 dogs and several cats are in an image
  - >> The classifier predicted the presence of 8 dogs; it properly identified 5 dogs as dogs (true positive), but it incorrectly identified 3 cats as dogs (false positive)
  - >> Precision =  $TP / (FP + TP) = 5 / (3 + 5) = 0.625$
  - >> Recall =  $5 / 12 = 0.4167$



Average Precision (AP) is the area under the precision-recall curve

mean Average Precision (AP) is the average of the AP of all classes. It provides an indication of how well objects are localized and identified.



How many selected items are relevant?

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

[https://en.wikipedia.org/wiki/Precision\\_and\\_recall](https://en.wikipedia.org/wiki/Precision_and_recall)


# Intersection over Union (1/2)



[https://www.aiukraine.com/wp-content/uploads/2016/09/Pashchenko\\_AIUkraine\\_2016.pdf](https://www.aiukraine.com/wp-content/uploads/2016/09/Pashchenko_AIUkraine_2016.pdf)


$$\text{Score} = \frac{\text{Area of overlap}}{\text{Area of union}}$$

*When the score exceeds a threshold (e.g. 0.5), and the object is correctly identified, a true positive is declared*




# Intersection over Union (2/2)

*IoU: 0.4034*




**Poor**

*IoU: 0.7330*



**Good**

*IoU: 0.9264*



**Excellent**

<https://www.pyimagesearch.com/2016/11/07/intersection-over-union-iou-for-object-detection/>



# mAP and IoU

- > The “goodness” of object recognition engines is evaluated using the following criteria
  - >> Was the *position* of the objects correctly identified => IoU
  - >> Were the objects correctly identified => mAP
- > When only the mAP value is given, the IoU is implicit in the “rulebook”
  - >> The [PASCAL VOC challenge](#) recognizes a true positive when IoU > 0.5
- > Engines using the [COCO](#) dataset usually show the IoU value used in the mAP calculation (e.g. mAP<sub>IoU</sub>, mAP@IoU)

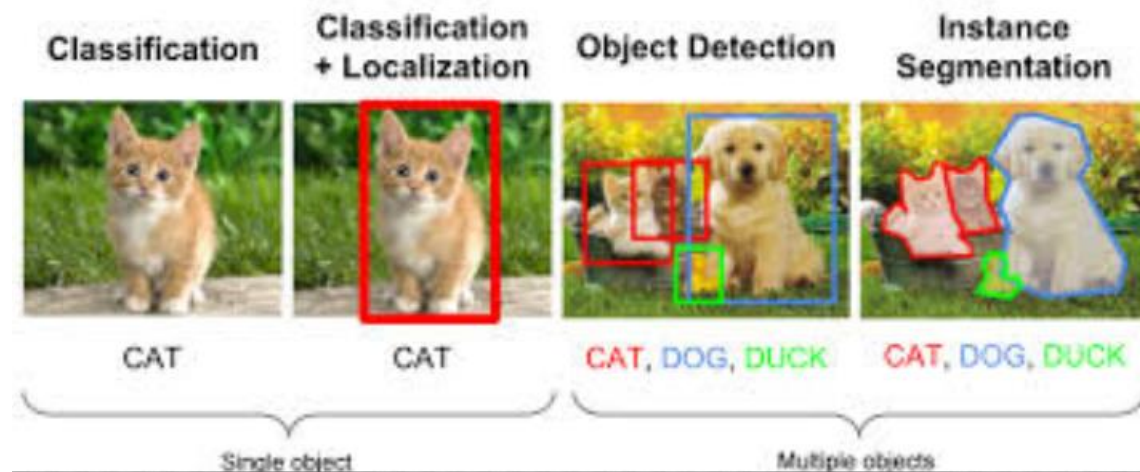


# Detection Background Overview



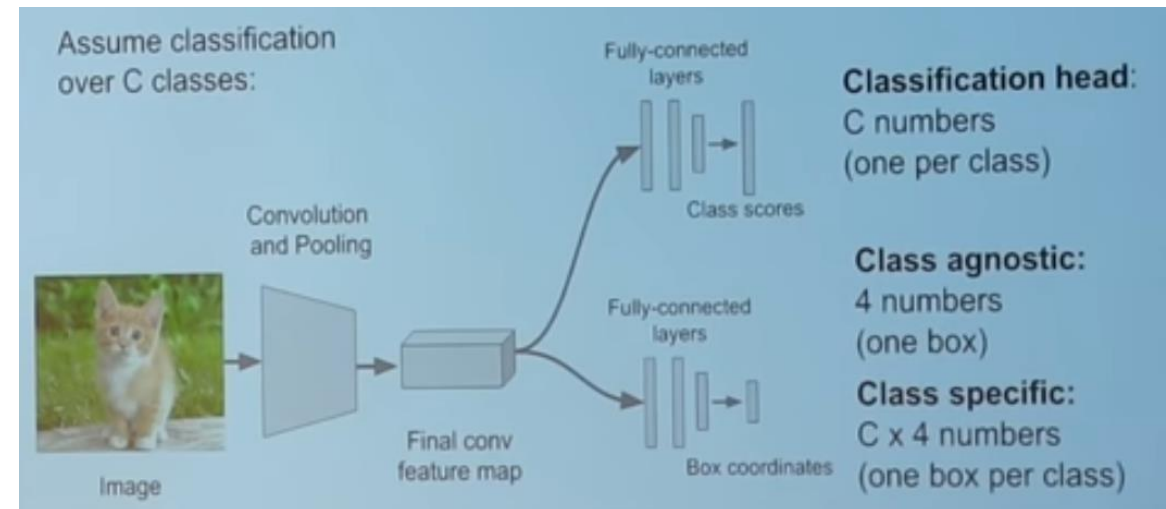
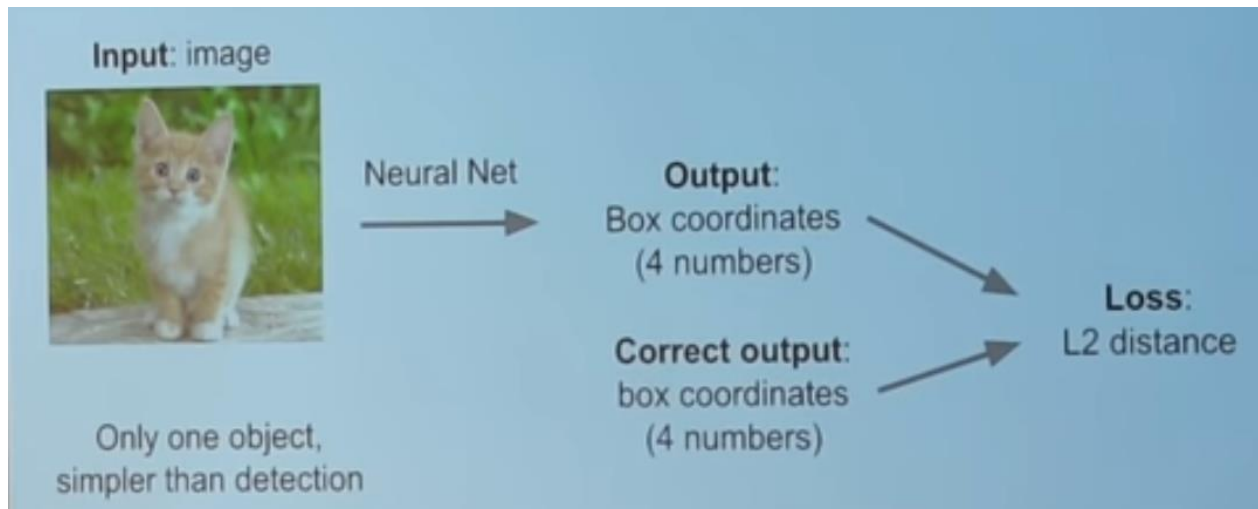
# Detection Background Information (1)

- > A great overview to start with for Detection networks (which doesn't include SSD) is available here: <https://youtu.be/2xtx-gk3PqY>
  - >> The following background information is a high level summary and the reference above does a much better job giving a detailed overview of the task at hand
- > **Assumption: we have already reviewed classification which produces a class output based on an input image**
  - The next step toward detection is localization
- > **Localization not only classifies the image but locates the object's instance within that image**



# Detection Background Information (2)

- > **Localization alone can be modeled as a regression problem which outputs (and can be trained on) a series of values (i.e. box coordinates) rather than only the distinct class**
- > **In addition to using the fully connected layers which are used in classification networks, a regression head is also used which will produce the box coordinates**
  - >> The image comes into the network and produces the class numbers as well as the box coordinates which come from the regression head
  - >> The correct output is also a 4 number value which can be trained against
    - The loss is then computed to train the network against these values
  - >> The Regression head can be attached after the convolution layers, or could also be attached after the last FC layer



# Detection Background Information (3)


- > **The localization technique can be extended for localizing more than one object**
  - >> Sliding window approaches can be used
    - Use a sliding window to identify classes/bounding boxes for sub-regions of an image, then combine them into a higher level score
- > **HOWEVER – for a detection problem, the outputs are variable**
  - >> need to detect an unknown number objects of different classes, varying scales, and potential overlapping regions in a single image
- > **One initial approach to address this issue was to model the detection problem as classification rather than regression**
  - >> Problem – need to test many different positions and scales
    - OLD Solution – if your classifier is fast enough you can just do it multiple times with sliding windows
    - R-CNN type Solution – region proposals could be used to act as a class agnostic object detector to look for blob-like regions
      - These types of networks are called “Region Proposal Networks”
      - This is complicated for training and slower – improvements were made in Fast R-CNN, and Faster R-CNN, but still lacks real time performance
        - Part of the reason is that two completely separate networks are needed – one for generating the region proposals, and one for the classification

# Detection Background Information (4)


- > **Even after many improvements to Region Proposal Networks (Fast R-CNN, and Faster R-CNN) the performance of these is still not “real time”**
  - >> ~5-7 FPS for Faster R-CNN on an x86+GPU machine (results are good, however in terms of mAP%)
- > **YOLO then introduced the idea of solving the detection problem as a regression problem**
  - >> This eliminates the need for the region proposal stage and greatly decreases processing time
  - >> This is done by discretizing the output space of bounding boxes into a set of default boxes over different aspect ratios and scales per feature map location
- > **YOLO uses a specific set of convolutional feature layers (i.e. backbone) called DarkNet**
  - >> Also trained in the darknet framework
  - >> SSD is similar to YOLO in that it is a single shot detector, but differs in a couple major ways
    - SSD can use feature extraction layers from many various classification networks (VGG, Inception, ResNet, etc.)
      - Feature layers extract edges, shapes, etc. of interest from the input data
    - SSD also uses a different strategy for the regression head which identifies the bounding box coordinates

# SSD Introduction


SSD is an unified framework for object detection with a single network. You can use the code to train/evaluate a network for object detection task. For more details, please refer to our [arXiv paper](#) and our [slide](#).



(a) Image with GT boxes



(b)  $8 \times 8$  feature map



(c)  $4 \times 4$  feature map

System	VOC2007 test <u>mAP</u>	FPS (Titan X)	Number of Boxes	Input resolution
<a href="#">Faster R-CNN (VGG16)</a>	73.2	7	~6000	~1000 x 600
<a href="#">YOLO (customized)</a>	63.4	45	98	448 x 448
SSD300* (VGG16)	77.2	46	8732	300 x 300
SSD512* (VGG16)	<b>79.8</b>	19	24564	512 x 512

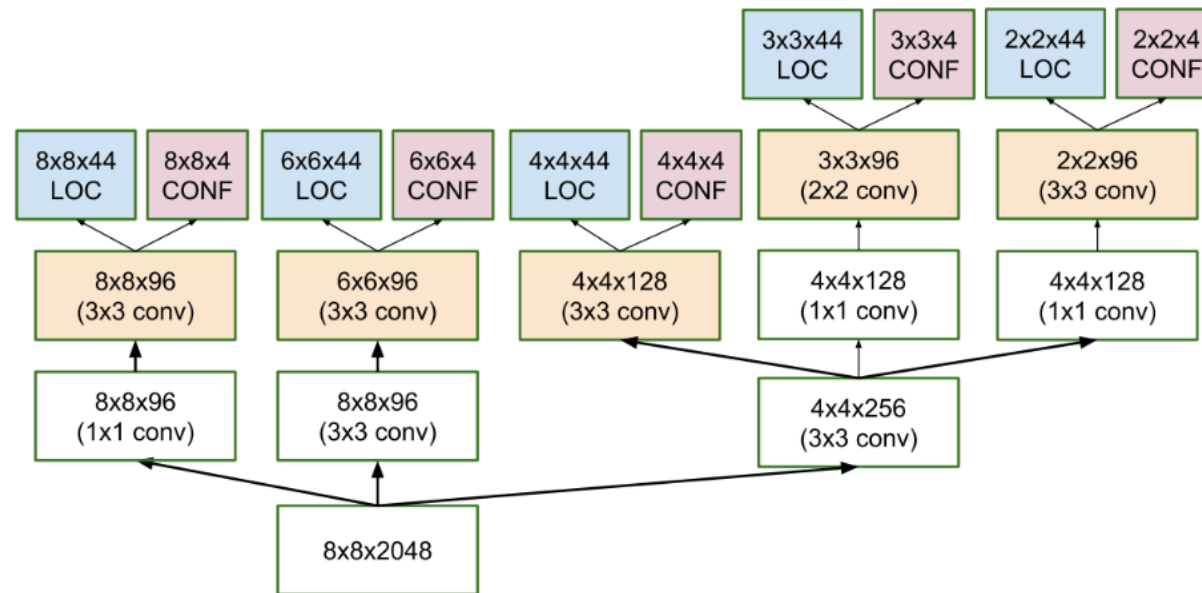
Method	VOC2007 test		VOC2012 test		COCO test-dev2015		
	07+12	07+12+COCO	07++12	07++12+COCO	trainval35k		
	0.5	0.5	0.5	0.5	0.5:0.95	0.5	0.75
SSD300*	77.2	81.2	75.8	79.3	25.1	43.1	25.8
SSD512*	<b>79.8</b>	<b>83.2</b>	<b>78.5</b>	<b>82.2</b>	<b>28.8</b>	<b>48.5</b>	<b>30.3</b>

Note: SSD300\* and SSD512\* are the latest models. Current code should reproduce these results.

Citation: <https://github.com/weiliu89/caffe/tree/ssd>

# MultiBox Background Information (1)

- > The Single Shot Detector (SSD) is based on the MultiBox work done by Szegedy
- > In MultiBox, the fundamental idea was to train a CNN that outputs coordinates of object boxes correctly and rank the proposals by their likelihood of being an accurate box for an object of interest
- > The primary outputs of the MultiBox Network are:
  - >> Confidence (CONF): logistic loss on the estimates of proposal corresponding to an object of interest
  - >> Location (LOC): loss corresponding to some similarity measurement between objects and closest matching object box prediction (uses L2 Norm)



Architecture of multi-scale convolutional prediction of the location and confidences of multibox


Source: <https://towardsdatascience.com/understanding-ssd-multibox-real-time-object-detection-in-deep-learning-495ef744fab>




# MultiBox Background Information (2)

- > **MultiBox is the portion of the network used as a reference for the bounding box regressor**
  - >> MultiBox starts with priors (or anchors – as described in R-CNN) which are pre-computed, fixed size bounding boxes
  - >> The initial predictions will then start with these prior boxes and training will attempt to regress them closer to the true bounding box locations
  - >> The images are divided into sub regions including 8x8, 6x6, 4x4 and the 1x1 is used to predict the single largest prior
    - Each region contains 11 potential boxes of differing aspect ratios/sizes


8x8 Grid



6x6 Grid

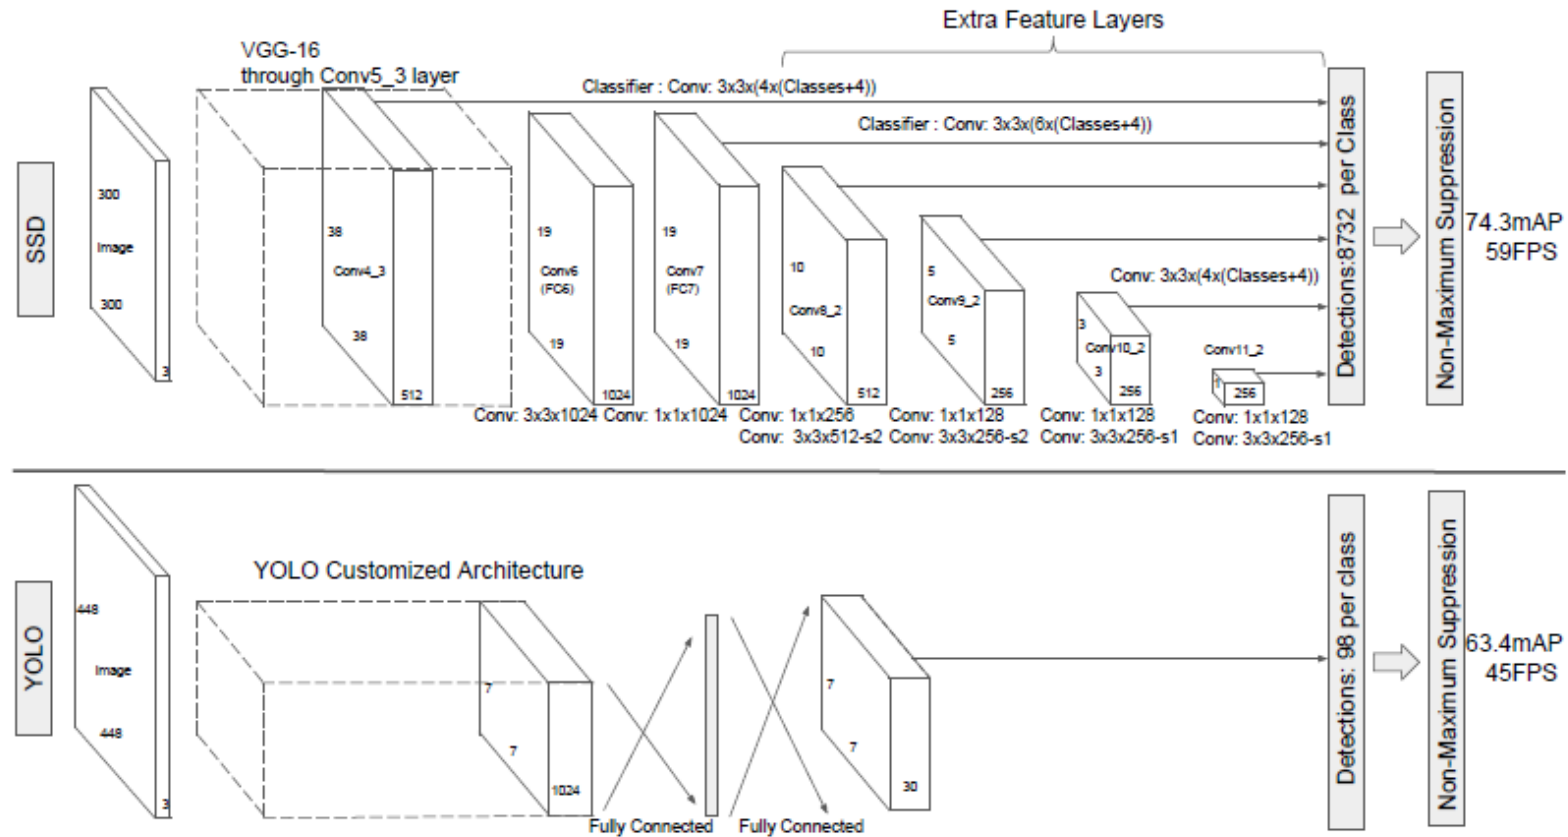


4x4 Grid



# SSD Background Information

- > SSD bounding box regressor is based on MultiBox and somewhat similar to YOLO
- > One of the key differences is that SSD uses convolutional layers to provide different image scales (stride=2 for downsampling)




Citation: <https://arxiv.org/abs/1512.02325>

# Regression Outputs from SSD

## > Rather than implement the exact Prior Boxes etc. from MultiBox, SSD implements the following:

- >> 38x38 grid with 4 boxes per cell = 5776 boxes
- >> 19x19 grid with 6 boxes per cell = 2166 boxes
- >> 10x10 grid with 6 boxes per cell = 600 boxes
- >> 5x5 grid with 6 boxes per cell = 150 boxes
- >> 3x3 grid with 4 boxes per cell = 36 boxes
- >> 1x1 grid with 4 boxes per cell = 4 boxes
- >> TOTAL 8732 boxes



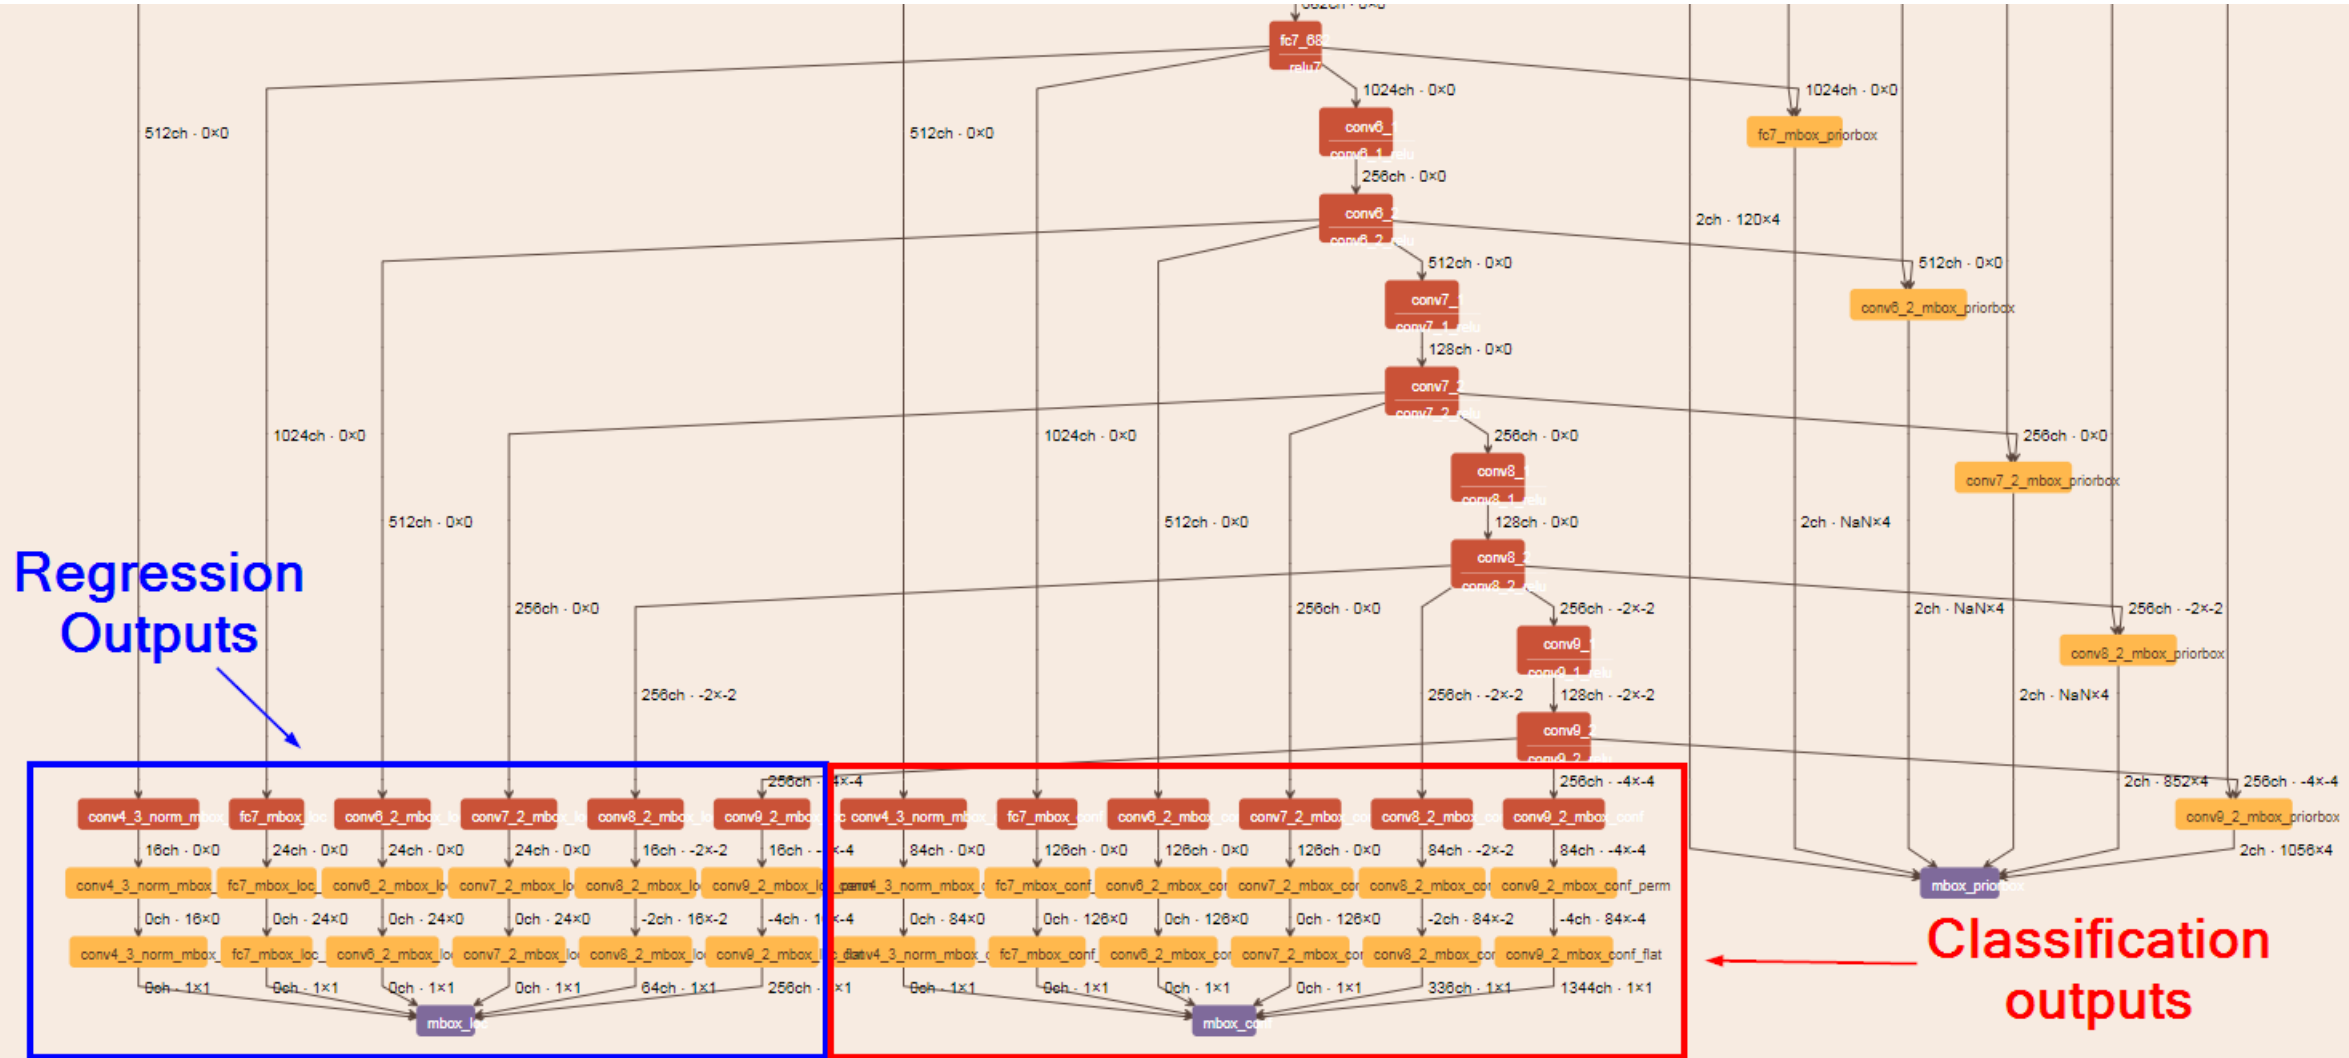
## > Convolution operations are used to create these outputs

- >> For Example:
  - 38x38 grid corresponds to a 38x38 feature map size which is the input to the convolution layer
    - 3x3 kernel size is used
  - Number of output channels (feature maps) from these layers depends on the number of boxes desired for that grid size
    - For 6 boxes, 24 output channels would be needed (4 values for each box)
    - For 4 boxes, 16 output channels would be needed (4 values for each box)

# Classification Outputs from SSD

- > For the classification outputs each of the grid sizes also will have a classification result
- > The number of output channels depends on the number of boxes per cell, grid size, and number of output classes
- > For PASCAL VOC, there are 21 output classes (20 objects and 1 background)
  - >> 38x38 grid -> 4 boxes/cell \* 21 classes = 84 output channels @ 38x38
  - >> 19x19 grid -> 6 boxes/cell \* 21 classes = 126 output channels @ 19x19
  - >> 10x10 grid -> 6 boxes/cell \* 21 classes = 126 output channels @ 10x10
  - >> 5x5 grid -> 6 boxes/cell \* 21 classes = 126 output channels @ 5x5
  - >> 3x3 grid -> 4 boxes/cell \* 21 classes = 84 output channels @ 3x3
  - >> 1x1 grid -> 4 boxes/cell \* 21 classes = 84 output channels @ 1x1
- > The output feature map size then relates back to the grid location and the channel relates to the class and specific priorbox for that class
- > These are finally run through softmax function

# SSD Classification/Regression Subset Netscope Diagram






# Detection Network Comparison




# Comparison of Detection Network Accuracy and Speed

- > **Warning** – these numbers shouldn't be taken as hard and fast comparison point as the experiments are done in different settings





# Performance Variations for Detection Framework and Feature Extractor



Source: [https://medium.com/@jonathan\\_hui/object-detection-speed-and-accuracy-comparison-faster-r-cnn-r-fcn-ssd-and-yolo-5425656ae359](https://medium.com/@jonathan_hui/object-detection-speed-and-accuracy-comparison-faster-r-cnn-r-fcn-ssd-and-yolo-5425656ae359)

# Detection Networks Comparisons

	SSD (300x300)	SSD (512x512)	YOLOv1	YOLOv2	Faster R-CNN
mAP % (VOC)	72.4	74.9	57.9	73.4	70.4
FPS (K40 GPU batch=1)	46	19	45	67	7
#Boxes	8732	24564	98		~6000
#Classes	20	20	20	Can be modified for Up to 9000	
Input Resolution	300x300	512x512	448x448	416x416	~1000x600

	SSD321	SSD513	YOLOv2	YOLOv3	R-FCN	DSSD321	DSSD513
mAP% (50 IoU - COCO)	45.4	50.4	44.0	55.3	51.9	46.1	53.3
FPS	16.3	8	45	34.5	11.7	11.7	6.4

Source: [https://medium.com/@jonathan\\_hui/object-detection-speed-and-accuracy-comparison-faster-r-cnn-r-fcn-ssd-and-yolo-5425656ae359](https://medium.com/@jonathan_hui/object-detection-speed-and-accuracy-comparison-faster-r-cnn-r-fcn-ssd-and-yolo-5425656ae359)

# Detection Networks Summary

- > RCNN/Fast R-CNN/Faster R-CNN is generally the highest overall for detection accuracies, but much slower in terms of performance
- > SSD is particularly good at detecting large objects, but not as good with smaller objects or objects that are close together
  - >> Different variants of SSD perform this task better than others
- > YOLO has variants that can detect up to 9000 object classes and YOLOv3 claims to be faster and more accurate than SSD (of course likely feature layer dependent)
  - >> YOLO encapsulates its own feature extraction layers (Darknet) whereas SSD and Faster R-CNN etc. are designed to work with other feature extractors such as VGG, Resnet, etc.
  - >> YOLOv3 makes improvements to YOLOv2 specifically for smaller objects

[Video Comparison of YOLOv2, YOLO9000, SSD, and Faster R-CNN](#)

[Yolov3 Video Demo](#)

**Adaptable.**  
**Intelligent.**

